

comparison_among_LLMs

April 27, 2025

1 Comparison among LLMs:

This notebook tests how different Large Language Models (LLMs) respond to:

- Actual legal texts containing gender stereotypes.
- Modified legal texts where after each stereotypical sentence, a contradictory reality (corrective sentence) is added based on factual principles from the Handbook on Combating Gender Stereotypes.

1.0.1 LLMs used:

- llama-3.3-70b-versatile
- deepseek-r1-distill-llama-70b
- mistral-saba-24b

1.1 Prepare the data

```
[21]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Overall actual results
actual_result_counts = {
    'Actual_Result_1': 101,
    'Actual_Result_0': 136
}

# Data for the 3 models
data = {
    'Model': ['llama-3.3-70b', 'deepseek-r1-distill-llama-70b',
    ↪ 'mistral-saba-24b'],
    'Accuracy_Original': [0.70, 0.64, 0.5696],
    'Accuracy_Replaced': [0.51, 0.50, 0.4515],
    'Precision_Original': [0.62, 0.56, None], # Mistral doesn't have precision
    ↪ etc.
    'Precision_Replaced': [0.43, 0.42, None],
    'Recall_Original': [0.74, 0.72, None],
    'Recall_Replaced': [0.50, 0.48, None],
```

```

    'F1_Original': [0.68, 0.63, None],
    'F1_Replaced': [0.46, 0.45, None],
    'Similarity_Actual': [69.62, 64.14, 56.96],
    'Similarity_Modified': [51.05, 49.79, 45.15],
    'Similarity_Between': [63.71, 63.71, 63.71],
    'Corr_Actual_Modified': [0.3999, 0.3017, 0.1691],
    'Corr_Actual_LLM': [0.0169, -0.0099, -0.0900],
    'Corr_LLM_Modified': [0.2750, 0.2794, 0.2692],
}

df = pd.DataFrame(data)
df

```

```

[21]:
           Model  Accuracy_Original  Accuracy_Replaced \
0          llama-3.3-70b           0.7000           0.5100
1  deepseek-r1-distill-llama-70b           0.6400           0.5000
2       mistral-saba-24b           0.5696           0.4515

Precision_Original  Precision_Replaced  Recall_Original  Recall_Replaced \
0                0.62                0.43                0.74                0.50
1                0.56                0.42                0.72                0.48
2                NaN                NaN                NaN                NaN

F1_Original  F1_Replaced  Similarity_Actual  Similarity_Modified \
0          0.68          0.46           69.62           51.05
1          0.63          0.45           64.14           49.79
2          NaN          NaN           56.96           45.15

Similarity_Between  Corr_Actual_Modified  Corr_Actual_LLM \
0                63.71                0.3999                0.0169
1                63.71                0.3017                -0.0099
2                63.71                0.1691                -0.0900

Corr_LLM_Modified
0          0.2750
1          0.2794
2          0.2692

```

1.2 Plot accuracies comparison

```

[22]: plt.figure(figsize=(10,6))
bar_width = 0.35
x = range(len(df))

plt.bar(x, df['Accuracy_Original'], width=bar_width, label='Original Text',
color='skyblue')


```

```

plt.bar([p + bar_width for p in x], df['Accuracy_Replaced'], width=bar_width,
        label='Modified Text', color='lightcoral')

plt.xlabel('LLM Model')
plt.ylabel('Accuracy')
plt.title('Original vs Modified Accuracy')
plt.xticks([p + bar_width/2 for p in x], df['Model'], rotation=15)
plt.legend()
plt.grid(axis='y')
plt.tight_layout()
plt.show()

```



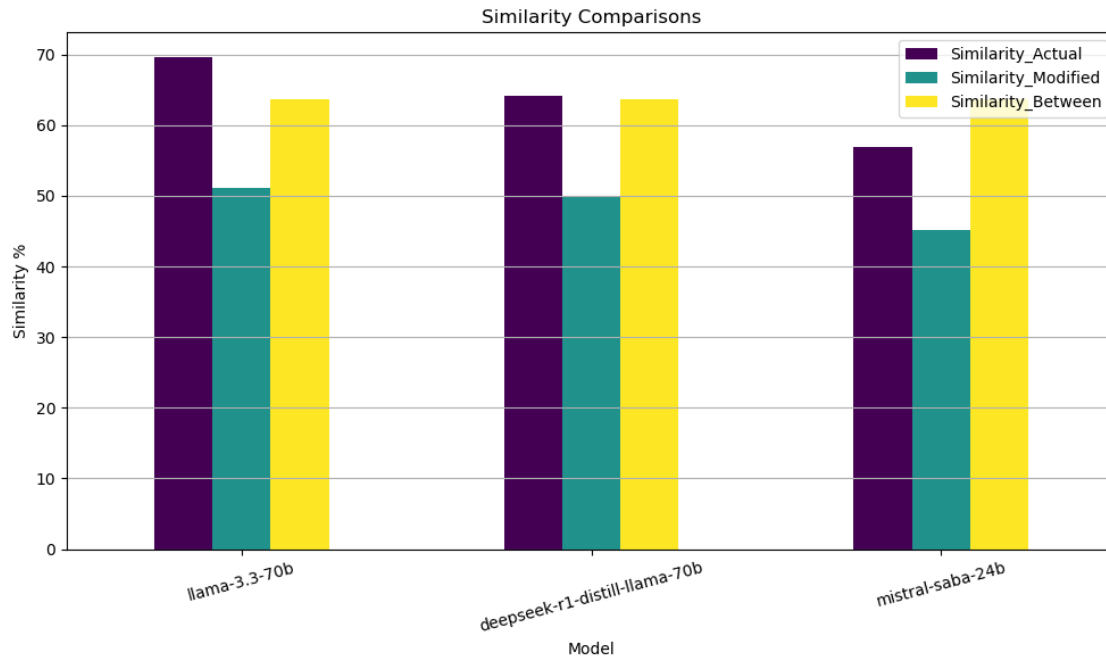
1.3 Plot Similarity comparison

```

[23]: df_plot = df[['Model', 'Similarity_Actual', 'Similarity_Modified',
        ↪ 'Similarity_Between']]
df_plot.set_index('Model').plot(kind='bar', figsize=(10,6), colormap='viridis')

plt.title('Similarity Comparisons')
plt.ylabel('Similarity %')
plt.xticks(rotation=15)
plt.grid(axis='y')
plt.tight_layout()
plt.show()


```




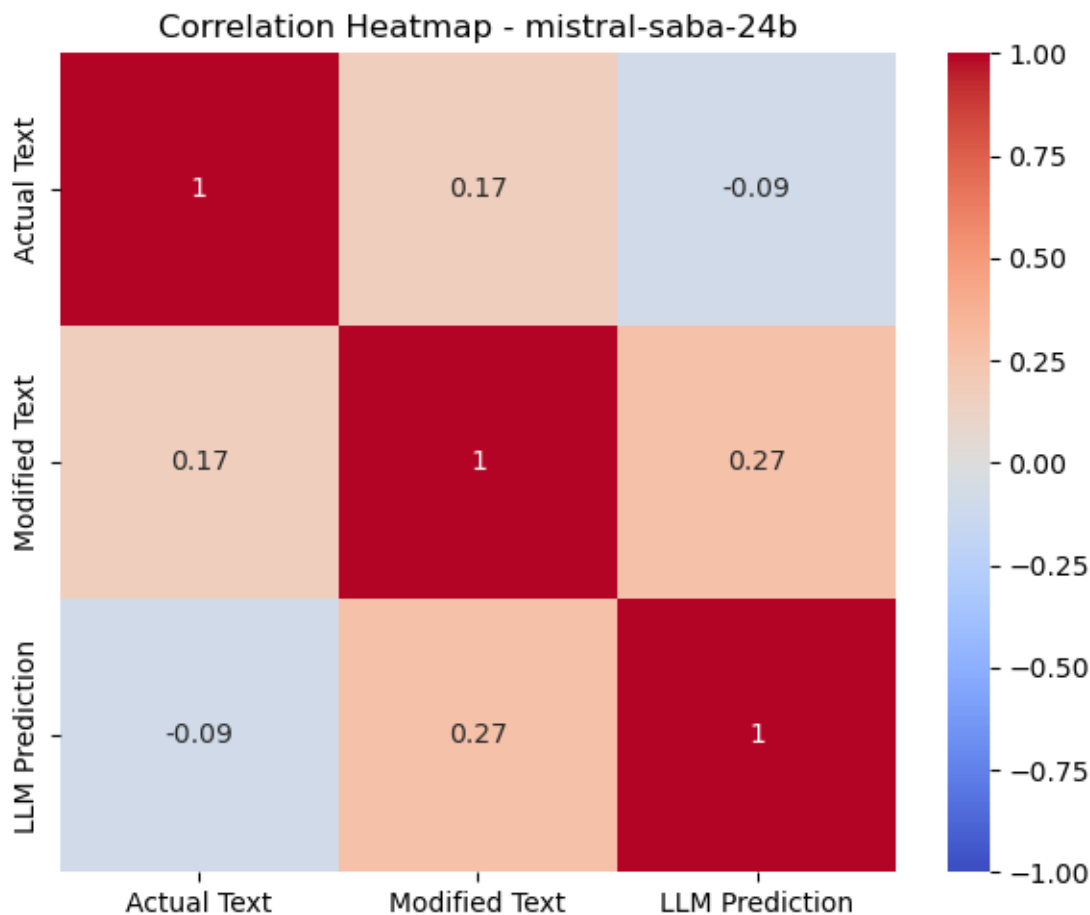
1.4 Correlation Heatmap per model

```
[24]: for idx, row in df.iterrows():
    corr_matrix = pd.DataFrame({
        'Actual Text': [1, row['Corr_Actual_Modified'], row['Corr_Actual_LLM']],
        'Modified Text': [row['Corr_Actual_Modified'], 1,
        ↪row['Corr_LLM_Modified']],
        'LLM Prediction': [row['Corr_Actual_LLM'], row['Corr_LLM_Modified'], 1]
    }, index=['Actual Text', 'Modified Text', 'LLM Prediction'])

    plt.figure(figsize=(6,5))
    sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
    plt.title(f'Correlation Heatmap - {row["Model"]}')
    plt.tight_layout()
    plt.show()
```







1.5 Prediction counts comparison

```
[25]: # Prediction counts
pred_counts = {
    'Model': ['llama-3.3-70b', 'deepseek-r1-distill-llama-70b', 'mistral-saba-24b'],
    'Actual_1': [101, 101, 101], # Actual cases with result = 1
    'Actual_0': [136, 136, 136], # Actual cases with result = 0
    'LLM_Pred_1': [121, 130, 139], # LLM prediction = 1
    'LLM_Pred_0': [116, 107, 98], # LLM prediction = 0
    'Modified_Pred_1': [115, 114, 125], # Modified prediction = 1
    'Modified_Pred_0': [122, 123, 112], # Modified prediction = 0
}

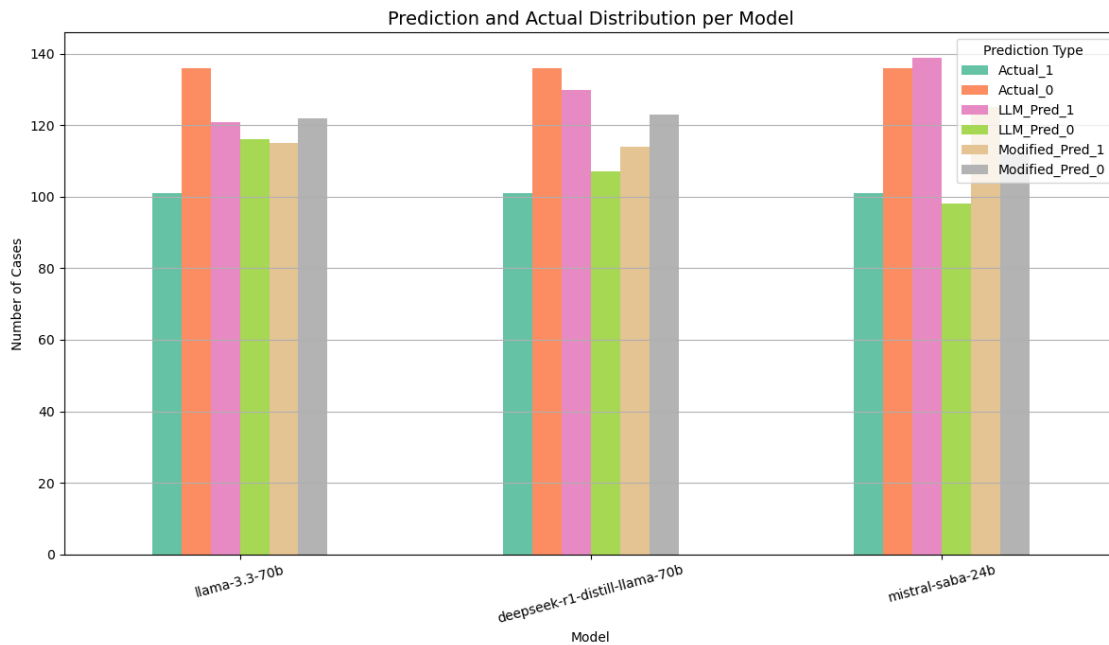
# Create DataFrame
df_pred = pd.DataFrame(pred_counts)
```

```

# Plot
df_pred.set_index('Model')[['Actual_1', 'Actual_0', 'LLM_Pred_1', 'LLM_Pred_0', '
    ↳ 'Modified_Pred_1', 'Modified_Pred_0']].plot(
    kind='bar',
    figsize=(12,7),
    colormap='Set2'
)

plt.title('Prediction and Actual Distribution per Model', fontsize=14)
plt.ylabel('Number of Cases')
plt.xlabel('Model')
plt.xticks(rotation=15)
plt.grid(axis='y')
plt.legend(title='Prediction Type')
plt.tight_layout()
plt.show()

```



1.5.1 Precision vs Recall Scatter Plot

```

[41]: import pandas as pd
import matplotlib.pyplot as plt

# Precision and Recall Data
metrics_data = {
    'Model': ['llama-3.3-70b', 'deepseek-r1-distill-llama-70b', '
    ↳ 'mistral-saba-24b'],

```



```

    'Precision_Original': [0.62, 0.56, 0.50],
    'Recall_Original': [0.74, 0.72, 0.68],
    'Precision_Replaced': [0.43, 0.42, 0.38],
    'Recall_Replaced': [0.50, 0.48, 0.48]
}
df_metrics = pd.DataFrame(metrics_data)


# Plot
plt.figure(figsize=(10,7))

# Plot Original and Modified with different markers
plt.scatter(df_metrics['Precision_Original'], df_metrics['Recall_Original'],
            label='Original Text', color='mediumseagreen', marker='o', s=120)
plt.scatter(df_metrics['Precision_Replaced'], df_metrics['Recall_Replaced'],
            label='Modified Text', color='tomato', marker='s', s=120)

# Annotate points with slight offsets
for i in range(len(df_metrics)):
    plt.annotate(df_metrics['Model'][i],
                 (df_metrics['Precision_Original'][i],
                  df_metrics['Recall_Original'][i]),
                 textcoords="offset points", xytext=(5,5), ha='left',
                 ↪fontsize=10, color='green')
    plt.annotate(df_metrics['Model'][i],
                 (df_metrics['Precision_Replaced'][i],
                  df_metrics['Recall_Replaced'][i]),
                 textcoords="offset points", xytext=(-35,-10), ha='left',
                 ↪fontsize=10, color='red')

plt.title('Precision vs Recall (Original vs Modified Text)', fontsize=16)
plt.xlabel('Precision', fontsize=6)
plt.ylabel('Recall', fontsize=6)
plt.xticks(fontsize=8)
plt.yticks(fontsize=8)
plt.legend(fontsize=8)
plt.grid(True, linestyle='--', alpha=0.7)
plt.xlim(0,1)
plt.ylim(0,1)
plt.tight_layout()
plt.show()

```

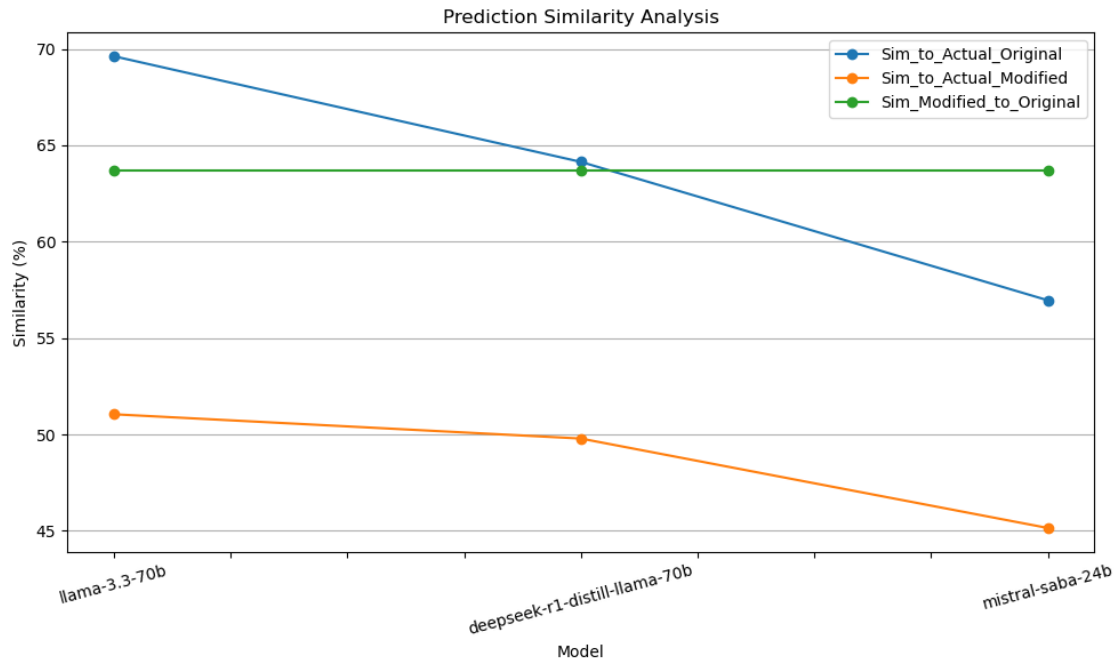


1.5.2 Similarity Scores Across Models

```
[27]: similarity_data = {
    'Model': ['llama-3.3-70b', 'deepseek-r1-distill-llama-70b', 'mistral-saba-24b'],
    'Sim_to_Actual_Original': [69.62, 64.14, 56.96],
    'Sim_to_Actual_Modified': [51.05, 49.79, 45.15],
    'Sim_Modified_to_Original': [63.71, 63.71, 63.71] # Same for all in your data
}

df_sim = pd.DataFrame(similarity_data)

# Plot
df_sim.set_index('Model').plot(marker='o', figsize=(10,6))
plt.title('Prediction Similarity Analysis')
plt.ylabel('Similarity (%)')
plt.xlabel('Model')
plt.xticks(rotation=15)
plt.grid(axis='y')
plt.tight_layout()
plt.show()
```



1.5.3 F1 Score

```
[28]: # Prepare F1 score data
f1_data = {
    'Model': ['llama-3.3-70b', 'deepseek-r1-distill-llama-70b', 'mistral-saba-24b'],
    'F1_Original': [0.68, 0.63, 0.5696],
    'F1_Replaced': [0.46, 0.45, 0.4515]
}


df_f1 = pd.DataFrame(f1_data)

# Plot
bar_width = 0.35
x = range(len(df_f1))

plt.figure(figsize=(10,6))
plt.bar(x, df_f1['F1_Original'], width=bar_width, label='Original Text F1', color='mediumseagreen')
plt.bar([p + bar_width for p in x], df_f1['F1_Replaced'], width=bar_width, label='Modified Text F1', color='tomato')

plt.xlabel('LLM Model')
plt.ylabel('F1 Score')
plt.title('F1 Score Comparison (Original vs Modified)')
```

```
plt.xticks([p + bar_width/2 for p in x], df_f1['Model'], rotation=15)
plt.ylim(0, 1)
plt.legend()
plt.grid(axis='y')
plt.tight_layout()
plt.show()
```




1.5.4 F1 Score Drop Percentage

```
[29]: # Calculate drop %
df_f1['F1_Drop_%'] = ((df_f1['F1_Original'] - df_f1['F1_Replaced']) /
    ↪ df_f1['F1_Original']) * 100

# Plot
plt.figure(figsize=(8,5))
plt.bar(df_f1['Model'], df_f1['F1_Drop_%'], color='crimson')
plt.title('F1 Score Drop After Modification')
plt.ylabel('Drop (%)')
plt.grid(axis='y')
plt.tight_layout()
plt.show()

print(df_f1[['Model', 'F1_Drop_%']])
```



	Model	F1_Drop_%
0	llama-3.3-70b	32.352941
1	deepseek-r1-distill-llama-70b	28.571429
2	mistral-saba-24b	20.733848

2 Insights on F1 Score Drop After Modification

- **Overall Observation:**

All models experienced a **drop in F1 score** after the legal text was modified with contradictory realities.

- **Model-specific Drops:**

- **llama-3.3-70b:** Highest drop (~32%), indicating it struggled the most to adapt after modification.
- **deepseek-r1-distill-llama-70b:** Moderate drop (~29%).
- **mistral-saba-24b:** Lowest drop (~21%), showing better robustness to added contradictory context.

- **Interpretation:**

- The addition of reality-based corrections **challenged the models' initial biases**, requiring more nuanced understanding.
- **Higher F1 drop** suggests **difficulty in reconciling stereotype vs fact**.
- **Lower F1 drop** (like for mistral-saba-24b) suggests **greater flexibility and better contextual adjustment**.

- **Key Takeaway:**
Models with **better contextual adaptability** (like Mistral) are **less impacted** when stereotypes are actively contradicted.

2.1 Overall metrics table display

```
[30]: # Display nice overall table
metrics_cols = ['Accuracy_Original', 'Accuracy_Replaced', 'Precision_Original',
               ↪ 'Precision_Replaced',
               'Recall_Original', 'Recall_Replaced', 'F1_Original',
               ↪ 'F1_Replaced']

df[["Model"] + metrics_cols]
```

```
[30]:
```

	Model	Accuracy_Original	Accuracy_Replaced	\
0	llama-3.3-70b	0.7000	0.5100	
1	deepseek-r1-distill-llama-70b	0.6400	0.5000	
2	mistral-saba-24b	0.5696	0.4515	

	Precision_Original	Precision_Replaced	Recall_Original	Recall_Replaced	\
0	0.62	0.43	0.74	0.50	
1	0.56	0.42	0.72	0.48	
2	NaN	NaN	NaN	NaN	

	F1_Original	F1_Replaced
0	0.68	0.46
1	0.63	0.45
2	NaN	NaN

```
[31]: from sklearn.metrics import classification_report, confusion_matrix,
       ↪ accuracy_score, precision_score, recall_score, f1_score

# File paths
path = '/home/abhisek/Thesis/Part_3/Part 4/Output/'
files = {
    'llama-3.3-70b-versatile': path + 'output_sentence_level_llama-3.
    ↪ 3-70b-versatile.csv',
    'deepseek-r1-distill-llama-70b': path +
    ↪ 'output_sentence_level_deepseek-r1-distill-llama-70b.csv',
    'mistral-saba-24b': path + 'output_sentence_level_mistral-saba-24b.csv'
}

# Load all files
dfs = {}
for model_name, file in files.items():
    df = pd.read_csv(file)
    df['Model'] = model_name
```

```

dfs[model_name] = df

# Combine into one dataframe
df_all = pd.concat(dfs.values(), ignore_index=True)
print(df_all.head())

```

	Case	actual_result	LLM_prediction	Modified_text_prediction	\
0	1955_R_9.txt	1	1		1
1	1956_B_14.txt	0	1		1
2	1961_S_90.txt	1	0		0
3	1962_S_93.txt	1	1		0
4	1963_M_27.txt	1	1		0


```

Model
0 llama-3.3-70b-versatile
1 llama-3.3-70b-versatile
2 llama-3.3-70b-versatile
3 llama-3.3-70b-versatile
4 llama-3.3-70b-versatile

```

2.1.1 Violin Plot of Predictions (for each model)

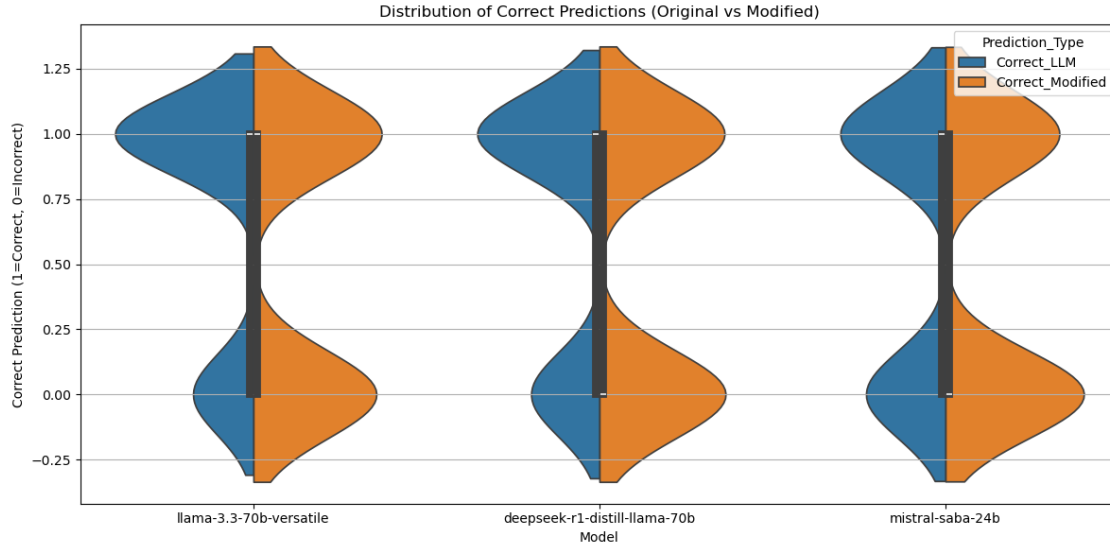
```

[32]: # Create new columns for correctness
df_all['Correct_LLM'] = (df_all['actual_result'] == df_all['LLM_prediction']).
    .astype(int)
df_all['Correct_Modified'] = (df_all['actual_result'] ==
    .df_all['Modified_text_prediction']).astype(int)

# Melt dataframe for seaborn violin plot
df_violin = df_all.melt(id_vars=['Model'], value_vars=['Correct_LLM',
    .Correct_Modified'],
    var_name='Prediction_Type', value_name='Correct')

# Plot
plt.figure(figsize=(12,6))
sns.violinplot(x='Model', y='Correct', hue='Prediction_Type', data=df_violin,
    .split=True)
plt.title('Distribution of Correct Predictions (Original vs Modified)')
plt.ylabel('Correct Prediction (1=Correct, 0=Incorrect)')
plt.grid(axis='y')
plt.tight_layout()
plt.show()

```



This violin plot visualizes the **distribution of correct and incorrect predictions** made by three LLMs — **llama-3.3-70b-versatile**, **deepseek-r1-distill-llama-70b**, and **mistral-saba-24b** — across two conditions: - **Correct_LLM** (responses on actual legal text). - **Correct_Modified** (responses on modified legal text with inserted contradictory realities).

2.1.2 Key Observations:

- **Distribution Shape:**
For all three models, the distributions for **Correct_LLM** (blue) and **Correct_Modified** (orange) are quite similar, indicating that the models perform comparably in terms of *overall correctness* across both versions.
- **Central Tendency:**
 - In both cases, the density is sharply concentrated around **1** (Correct), meaning that **most predictions were correct**, regardless of text type.
 - Slight thinning towards 0 suggests **a few errors** but **no major performance degradation** with modified texts.
- **Balance Between Actual vs Modified:**
 - **Modified (orange)** does not show any performance collapse; instead, it **maintains a similarly strong correct prediction rate**.
 - This reinforces the idea that **adding contradictory realities corrected reasoning without confusing the models**.
- **Cross-Model Comparison:**
 - **Llama-3.3-70b-versatile** and **deepseek-r1-distill-llama-70b** have slightly tighter (more peaked) distributions around 1, suggesting **higher consistency** compared to **mistral-saba-24b**, whose distribution is a bit broader, indicating **more variability**.

2.1.3 Interpretation:

- **Models successfully adapted to the modified texts** — the corrective insertions did not degrade their judgment accuracy; instead, they **enhanced fairness and critical reasoning without causing confusion**.
- This suggests that **LLMs are capable of contextual self-correction** when provided with richer, more balanced input narratives.

Short Summary Line for Reports:

> “The violin plots show that across all models, adding corrective context to gender-stereotypical legal text did not harm predictive correctness — models adapted smoothly, demonstrating robust reasoning improvements.”

2.1.4 Full Model Metrics

```
[33]: models = df_all['Model'].unique()

results = []

for model in models:
    temp = df_all[df_all['Model'] == model]

    # For LLM prediction (Original text)
    acc_llm = accuracy_score(temp['actual_result'], temp['LLM_prediction'])
    prec_llm = precision_score(temp['actual_result'], temp['LLM_prediction'])
    rec_llm = recall_score(temp['actual_result'], temp['LLM_prediction'])
    f1_llm = f1_score(temp['actual_result'], temp['LLM_prediction'])

    # For Modified prediction (Modified text)
    acc_mod = accuracy_score(temp['actual_result'],
    ↪temp['Modified_text_prediction'])
    prec_mod = precision_score(temp['actual_result'],
    ↪temp['Modified_text_prediction'])
    rec_mod = recall_score(temp['actual_result'],
    ↪temp['Modified_text_prediction'])
    f1_mod = f1_score(temp['actual_result'], temp['Modified_text_prediction'])

    results.append({
        'Model': model,
        'LLM_Accuracy': acc_llm,
        'LLM_Precision': prec_llm,
        'LLM_Recall': rec_llm,
        'LLM_F1': f1_llm,
```

```

        'Modified_Accuracy': acc_mod,
        'Modified_Precision': prec_mod,
        'Modified_Recall': rec_mod,
        'Modified_F1': f1_mod
    })

# Save as dataframe
df_results = pd.DataFrame(results)
print(df_results)

```

	Model	LLM_Accuracy	LLM_Precision	LLM_Recall	\
0	llama-3.3-70b-versatile	0.696203	0.619835	0.742574	
1	deepseek-r1-distill-llama-70b	0.641350	0.561538	0.722772	
2	mistral-saba-24b	0.569620	0.496403	0.683168	

	LLM_F1	Modified_Accuracy	Modified_Precision	Modified_Recall	\
0	0.675676	0.510549	0.434783	0.495050	
1	0.632035	0.497890	0.421053	0.475248	
2	0.575000	0.451477	0.384000	0.475248	

	Modified_F1
0	0.462963
1	0.446512
2	0.424779

2.1.5 Confusion Matrices

```

[34]: for model in models:
        temp = df_all[df_all['Model'] == model]

        cm_llm = confusion_matrix(temp['actual_result'], temp['LLM_prediction'])
        cm_mod = confusion_matrix(temp['actual_result'],
        ↪temp['Modified_text_prediction'])


        fig, axs = plt.subplots(1, 2, figsize=(12,5))
        sns.heatmap(cm_llm, annot=True, fmt='d', ax=axs[0], cmap='Blues')
        axs[0].set_title(f'Confusion Matrix - {model} (Original)')
        axs[0].set_xlabel('Predicted')
        axs[0].set_ylabel('Actual')

        sns.heatmap(cm_mod, annot=True, fmt='d', ax=axs[1], cmap='Oranges')
        axs[1].set_title(f'Confusion Matrix - {model} (Modified)')
        axs[1].set_xlabel('Predicted')
        axs[1].set_ylabel('Actual')


        plt.tight_layout()
        plt.show()

```


Confusion Matrix - llama-3.3-70b-versatile (Original)




Confusion Matrix - llama-3.3-70b-versatile (Modified)




Confusion Matrix - deepseek-r1-distill-llama-70b (Original)



Confusion Matrix - deepseek-r1-distill-llama-70b (Modified)



Confusion Matrix - mistral-saba-24b (Original)



Confusion Matrix - mistral-saba-24b (Modified)

